

# Learning Conditional Random Field parameters from training data

A review

Przemyslaw Polewski

Photogrammetry and Remote Sensing  
Technische Universität München

Department of Geoinformatics  
Hochschule München

09.11.2016 / PF-Seminar



# Outline

- 1 Conditional Random Fields
- 2 Direct methods
  - Probabilistic modeling
  - Grid search
- 3 Maximum likelihood variations
  - Pseudo-likelihood
  - Marginal approximations
- 4 Max-margin learning
  - Structured SVM
- 5 Software libraries

# Conditional Random Fields

## Definition

- Consider set of objects  $O = \{o_1, \dots, o_n\}$  with features  $X = \{x_1, \dots, x_n\}$  and labels  $Y = \{y_1, \dots, y_n\}$ .

# Conditional Random Fields

## Definition

- Consider set of objects  $O = \{o_1, \dots, o_n\}$  with features  $X = \{x_1, \dots, x_n\}$  and labels  $Y = \{y_1, \dots, y_n\}$ .
- Let  $G$  be a discrete structure (graph, lattice) with a neighborhood relation  $\sim$  on the elements of  $O$

# Conditional Random Fields

## Definition

- Consider set of objects  $O = \{o_1, \dots, o_n\}$  with features  $X = \{x_1, \dots, x_n\}$  and labels  $Y = \{y_1, \dots, y_n\}$ .
- Let  $G$  be a discrete structure (graph, lattice) with a neighborhood relation  $\sim$  on the elements of  $O$
- CRF: Undirected probabilistic graphical model which describes *conditional* probability (over  $G$ ) of object labels given their features:

# Conditional Random Fields

## Definition

- Consider set of objects  $O = \{o_1, \dots, o_n\}$  with features  $X = \{x_1, \dots, x_n\}$  and labels  $Y = \{y_1, \dots, y_n\}$ .
- Let  $G$  be a discrete structure (graph, lattice) with a neighborhood relation  $\sim$  on the elements of  $O$
- CRF: Undirected probabilistic graphical model which describes *conditional* probability (over  $G$ ) of object labels given their features:

$$P(y|x) = \frac{1}{Z(x)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x)\right] \quad (1)$$

# Conditional Random Fields

## Definition

- Consider set of objects  $O = \{o_1, \dots, o_n\}$  with features  $X = \{x_1, \dots, x_n\}$  and labels  $Y = \{y_1, \dots, y_n\}$ .
- Let  $G$  be a discrete structure (graph, lattice) with a neighborhood relation  $\sim$  on the elements of  $O$
- CRF: Undirected probabilistic graphical model which describes *conditional* probability (over  $G$ ) of object labels given their features:

$$P(y|x) = \frac{1}{Z(x)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x)\right] \quad (1)$$

Equivalent energy formulation:

$$E(y) = -\log P(y|x) = \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x) - \log Z(x) \quad (2)$$

# Conditional Random Fields

## Notation

$$E(y; \overbrace{\Theta}^{\text{param.}}) = \sum_{\substack{c \in \mathcal{C}(G) \\ \underbrace{\hspace{1cm}}_{\text{cliques of } G}}} \overbrace{\psi_c(y_c|x; \Theta)}^{\text{clique potentials}} - \log \underbrace{Z(x; \Theta)}_{\text{partition function}}$$



# Conditional Random Fields

## Notation

$$E(y; \overbrace{\Theta}^{\text{param.}}) = \sum_{c \in \underbrace{\mathcal{C}(G)}_{\text{cliques of } G}} \overbrace{\psi_c(y_c|x; \Theta)}^{\text{clique potentials}} - \log \underbrace{Z(x; \Theta)}_{\text{partition function}}$$



- $\Theta$ : vector of parameters  $\theta_1, \theta_2, \dots, \theta_m$
- $Z$ : function of data and parameters, but not of labels

# Conditional Random Fields

## Notation

$$E(y; \overbrace{\Theta}^{\text{param.}}) = \sum_{c \in \underbrace{\mathcal{C}(G)}_{\text{cliques of } G}} \overbrace{\psi_c(y_c|x; \Theta)}^{\text{clique potentials}} - \log \underbrace{Z(x; \Theta)}_{\text{partition function}}$$

- $\Theta$ : vector of parameters  $\theta_1, \theta_2, \dots, \theta_m$
- $Z$ : function of data and parameters, but not of labels

Two complementary tasks:

# Conditional Random Fields

## Notation

$$E(y; \overbrace{\Theta}^{\text{param.}}) = \sum_{\substack{c \in \mathcal{C}(G) \\ \underbrace{\hspace{1cm}}_{\text{cliques of } G}}} \overbrace{\psi_c(y_c|x; \Theta)}^{\text{clique potentials}} - \log \underbrace{Z(x; \Theta)}_{\text{partition function}}$$

- $\Theta$ : vector of parameters  $\theta_1, \theta_2, \dots, \theta_m$
- $Z$ : function of data and parameters, but not of labels

Two complementary tasks:

- Inference: minimize  $E$  w.r.t. **labels**  $y$  given parameters

# Conditional Random Fields

## Notation

$$E(y; \overbrace{\Theta}^{\text{param.}}) = \sum_{c \in \underbrace{\mathcal{C}(G)}_{\text{cliques of } G}} \overbrace{\psi_c(y_c|x; \Theta)}^{\text{clique potentials}} - \log \underbrace{Z(x; \Theta)}_{\text{partition function}}$$

- $\Theta$ : vector of parameters  $\theta_1, \theta_2, \dots, \theta_m$
- $Z$ : function of data and parameters, but not of labels



Two complementary tasks:

- Inference: minimize  $E$  w.r.t. **labels**  $y$  given parameters
- Learning: minimize  $E$  w.r.t. **parameters**  $\Theta$  given training labels

# Conditional Random Fields

## Neighborhood examples

Raster image


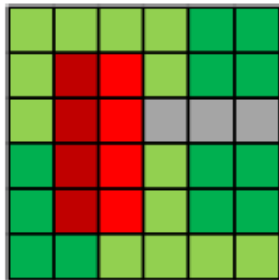


©Wegner, 2011

# Conditional Random Fields


## Neighborhood examples

Raster image



©Wegner, 2011

Point cloud



©pointclouds.org

# Literal interpretation of potentials

- Treat each clique potential as a classifier yielding a categorical distribution over possible class assignments to clique members

# Literal interpretation of potentials

- Treat each clique potential as a classifier yielding a categorical distribution over possible class assignments to clique members
- Each classifier can be trained independently using the user's ML method of choice (e.g. SVM, LR, RF)



# Literal interpretation of potentials

- Treat each clique potential as a classifier yielding a categorical distribution over possible class assignments to clique members
- Each classifier can be trained independently using the user's ML method of choice (e.g. SVM, LR, RF)
- Drawback: need a lot of training examples for rare class combinations

# Literal interpretation of potentials

- Treat each clique potential as a classifier yielding a categorical distribution over possible class assignments to clique members
- Each classifier can be trained independently using the user's ML method of choice (e.g. SVM, LR, RF)
- Drawback: need a lot of training examples for rare class combinations
- E.g. for unary and binary cliques:

$$E(y) = - \sum_i \log P(y_i | x) - \sum_i \sum_{j \in N(i)} \log P(\underbrace{y_i, y_j}_{\text{for all class pairs}} | x) \quad (3)$$

# Example

Classification of urban ALS point clouds [Niemeyer et al., 2014]

# Example

Classification of urban ALS point clouds [Niemeyer et al., 2014]

- 7 unbalanced object classes: grassland, road, building with gable roof, low vegetation, façade, building with flat roof, tree
- Unary and binary cliques

# Example

Classification of urban ALS point clouds [Niemeyer et al., 2014]

- 7 unbalanced object classes: grassland, road, building with gable roof, low vegetation, façade, building with flat roof, tree
- Unary and binary cliques



©Niemeyer et al. 2014

# Example

Classification of urban ALS point clouds [Niemeyer et al., 2014]

- 7 unbalanced object classes: grassland, road, building with gable roof, low vegetation, façade, building with flat roof, tree
- Unary and binary cliques

“A comparison of three different versions of a CRF-based classifier has shown that Random Forests are well suited for the computation of unary and pairwise potentials needed for CRFs(...)”

— Niemeyer et al.



©Niemeyer et al. 2014

## Grid search / empirical approach

- Applicable when number of parameters is small and their domain is known a priori better than  $(-\infty; \infty)$  (e.g.  $[0;1]$ )
- Grid search on parameter combinations  $\Theta^k \in \Theta$  with repeated inference:

## Grid search / empirical approach

- Applicable when number of parameters is small and their domain is known a priori better than  $(-\infty; \infty)$  (e.g.  $[0;1]$ )
- Grid search on parameter combinations  $\Theta^k \in \Theta$  with repeated inference:

---

```
1: for  $k = 1$  to  $|\Theta|$  do  
2:    $y^k = \arg \min_y E(y|x; \Theta^k)$   
3:    $e^k = \text{loss}(y^k, y^*)$   
4: end for  
5:  $k^* = \arg \min_k e^k$   
6: return  $\Theta^{k^*}$ 
```

---





## Grid search / empirical approach

- Applicable when number of **parameters** is small and their domain is known a priori better than  $(-\infty; \infty)$  (e.g.  $[0;1]$ )
- Grid search on parameter combinations  $\Theta^k \in \Theta$  with repeated inference:

---

```
1: for  $k = 1$  to  $|\Theta|$  do  
2:    $y^k = \arg \min_y E(y|x; \Theta^k)$   
3:    $e^k = \text{loss}(y^k, y^*)$   
4: end for  
5:  $k^* = \arg \min_k e^k$   
6: return  $\Theta^{k^*}$ 
```

---

Will (almost) never find the true optimal weights

- Discrete grid  $\neq$  continuous weights

# Empirical approach

## Example

Contrast sensitive Potts model for MLS point clouds [Weinmann et al., 2015]

# Empirical approach

## Example

Contrast sensitive Potts model for MLS point clouds [Weinmann et al., 2015]

- 5 unbalanced object classes: wire, pole/trunk, facade, ground, vegetation
- Unary cliques: learned with Random Forest
- Binary cliques: contrast-sensitive Potts model:

# Empirical approach

## Example

Contrast sensitive Potts model for MLS point clouds [Weinmann et al., 2015]

- 5 unbalanced object classes: wire, pole/trunk, facade, ground, vegetation
- Unary cliques: learned with Random Forest
- Binary cliques: contrast-sensitive Potts model:

$$\psi(y_i, y_j; x) = [y_i = y_j] \cdot \mathbf{w}_1 \cdot \frac{N_a}{N_{k_i}} \cdot [\mathbf{w}_2 + (1 - \mathbf{w}_2) \cdot e^{-\frac{d_{ij}(x)^2}{2\sigma^2}}]$$

# Empirical approach

## Example

Contrast sensitive Potts model for MLS point clouds [Weinmann et al., 2015]

- 5 unbalanced object classes: wire, pole/trunk, facade, ground, vegetation
- **Unary cliques**: learned with Random Forest
- **Binary cliques**: contrast-sensitive **Potts model**:

$$\psi(y_i, y_j; x) = [y_i = y_j] \cdot w_1 \cdot \frac{N_a}{N_{k_i}} \cdot [w_2 + (1 - w_2) \cdot e^{-\frac{d_{ij}(x)^2}{2\sigma^2}}]$$



“The weight parameters  $w_1$  and  $w_2$  could be set based on (...). Here, they are set to values that were found **empirically**.”

— Weinmann et al.

# Learning parameters with maximum likelihood

## The problem

Likelihood function:

$$\mathcal{L}(\Theta) = P(y|x; \Theta) = \frac{1}{Z(x, \Theta)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x; \Theta)\right]$$

# Learning parameters with maximum likelihood

## The problem

Likelihood function:

$$\mathcal{L}(\Theta) = P(y|x; \Theta) = \frac{1}{Z(x, \Theta)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x; \Theta)\right]$$

Partition “constant” is a *function* of parameters !

# Learning parameters with maximum likelihood

## The problem

Likelihood function:

$$\mathcal{L}(\Theta) = P(y|x; \Theta) = \frac{1}{Z(x, \Theta)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x; \Theta)\right]$$

Partition “constant” is a *function* of parameters !

$$Z(x, \Theta) = \underbrace{\sum_{y'}}_{\text{for all possible labelings}} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y'_c|x; \Theta)\right]$$



# Learning parameters with maximum likelihood

## The problem

Likelihood function:

$$\mathcal{L}(\Theta) = P(y|x; \Theta) = \frac{1}{Z(x, \Theta)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x; \Theta)\right]$$

Partition “constant” is a *function* of parameters !

$$Z(x, \Theta) = \underbrace{\sum_{y'} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y'_c|x; \Theta)\right]}_{\text{for all possible labelings}}$$

Exact computation infeasible even for small training sets and binary labeling...

# Learning parameters with maximum likelihood

## The problem

Likelihood function:

$$\mathcal{L}(\Theta) = P(y|x; \Theta) = \frac{1}{Z(x, \Theta)} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y_c|x; \Theta)\right]$$

Partition “constant” is a *function* of parameters !

$$Z(x, \Theta) = \underbrace{\sum_{y'} \exp\left[- \sum_{c \in \mathcal{C}(G)} \psi_c(y'_c|x; \Theta)\right]}_{\text{for all possible labelings}}$$

Exact computation infeasible even for small training sets and binary labeling...

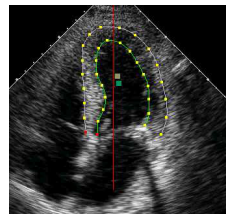
... **need an approximation** !

# Taming the partition function

## Method I

Case study: Heart Motion Abnormality Detection [Schmidt et al., 2008]

- Binary classification, up to pairwise cliques



©Schmidt et al. 2008

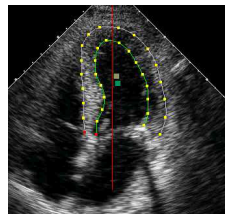
# Taming the partition function

## Method I

Case study: Heart Motion Abnormality Detection [Schmidt et al., 2008]

- Binary classification, up to pairwise cliques
- Assumes linear potentials:

$$\psi(y_i|x) = v^T F(x_i, y_i), \psi(y_i, y_j|x) = w^T F(x_{i,j}, y_i, y_j)$$



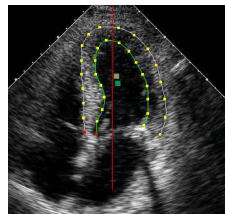
©Schmidt et al. 2008

# Taming the partition function

## Method I

Case study: Heart Motion Abnormality Detection [Schmidt et al., 2008]

- Binary classification, up to pairwise cliques
- Assumes linear potentials:  
 $\psi(y_i|x) = v^T F(x_i, y_i)$ ,  $\psi(y_i, y_j|x) = w^T F(x_{i,j}, y_i, y_j)$
- Compound parameter:  $\Theta = (v, w)$ , compound features  $F(x, y)$



©Schmidt et al. 2008

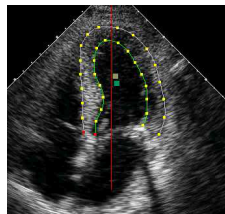
# Taming the partition function

## Method I

Case study: Heart Motion Abnormality Detection [Schmidt et al., 2008]

- Binary classification, up to pairwise cliques
- Assumes linear potentials:  
 $\psi(y_i|x) = v^T F(x_i, y_i)$ ,  $\psi(y_i, y_j|x) = w^T F(x_{i,j}, y_i, y_j)$
- Compound parameter:  $\Theta = (v, w)$ , compound features  $F(x, y)$

$$\ell(\Theta) = \log \mathcal{L}(\Theta) = \Theta^T F(x, y) - \log Z(x, \Theta)$$



©Schmidt et al. 2008

# Taming the partition function


## Method I

Case study: Heart Motion Abnormality Detection [Schmidt et al., 2008]

- Binary classification, up to pairwise cliques
- Assumes linear potentials:  
 $\psi(y_i|x) = v^T F(x_i, y_i)$ ,  $\psi(y_i, y_j|x) = w^T F(x_{i,j}, y_i, y_j)$
- Compound parameter:  $\Theta = (v, w)$ , compound features  $F(x, y)$

$$\ell(\Theta) = \log \mathcal{L}(\Theta) = \Theta^T F(x, y) - \log Z(x, \Theta)$$

$$\nabla \ell(\Theta) = F(x, y) - \underbrace{\sum_{y'} P(y'|x, \Theta) F(x, y')}_{\text{computationally infeasible}}$$



©Schmidt et al. 2008

# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:



# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:

Let  $Y = (Y_1, Y_2, \dots, Y_m)$  be a random vector with graph  $G = (V, E)$  defining conditional independence structure. Then:

# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:

Let  $Y = (Y_1, Y_2, \dots, Y_m)$  be a random vector with graph  $G = (V, E)$  defining conditional independence structure. Then:

$$P(y) \approx \prod_i P(y_i | y_k : (v_i, v_k) \in E)$$

# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:

Let  $Y = (Y_1, Y_2, \dots, Y_m)$  be a random vector with graph  $G = (V, E)$  defining conditional independence structure. Then:

$$P(y) \approx \prod_i P(y_i | y_k : (v_i, v_k) \in E)$$

Translating to our problem:

$$PL(\Theta) = \prod_i P(y_i | x, y_k : o_i \sim o_k; \Theta)$$

# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:

Let  $Y = (Y_1, Y_2, \dots, Y_m)$  be a random vector with graph  $G = (V, E)$  defining conditional independence structure. Then:

$$P(y) \approx \prod_i P(y_i | y_k : (v_i, v_k) \in E)$$

Translating to our problem:

$$PL(\Theta) = \prod_i P(y_i | x, y_k : o_i \sim o_k; \Theta) = \prod_i \exp(\Theta^T \overbrace{F_i(x, y)}^{\text{sub-features induced by } \sim_i}) / Z_i$$

# Taming the partition function

## Method I

Solution: use **pseudo-likelihood** approximation:

Let  $Y = (Y_1, Y_2, \dots, Y_m)$  be a random vector with graph  $G = (V, E)$  defining **conditional independence structure**. Then:

$$P(y) \approx \prod_i P(y_i | y_k : (v_i, v_k) \in E)$$

Translating to our problem:

$$PL(\Theta) = \prod_i P(y_i | x, y_k : o_i \sim o_k; \Theta) = \prod_i \exp(\Theta^T \underbrace{F_i(x, y)}_{\text{sub-features induced by } \sim_i}) / Z_i$$

$$Z_i = \sum_{\underbrace{y'_i}_{\text{only over current variable}}} \exp(\Theta^T F_i(x, y'_i))$$

only over current variable

# Taming the partition function

## Method II

Instead of approximating entire likelihood, try to estimate the problematic sum [Kumar et al., 2005]. Assume linear potentials ( $\Theta = (w, v)$ ):

# Taming the partition function

## Method II

Instead of approximating entire likelihood, try to estimate the problematic sum [Kumar et al., 2005]. Assume linear potentials ( $\Theta = (w, v)$ ):

$$\psi(y_i|x) = \log P(y_i|x) = y_i w^T x_i$$

$$\psi(y_i, y_j|x) = \log P(y_i, y_j|x) = y_i y_j v^T x_{i,j}$$

# Taming the partition function

## Method II

Instead of approximating entire likelihood, try to estimate the problematic sum [Kumar et al., 2005]. Assume linear potentials ( $\Theta = (w, v)$ ):

$$\psi(y_i|x) = \log P(y_i|x) = y_i w^T x_i$$

$$\psi(y_i, y_j|x) = \log P(y_i, y_j|x) = y_i y_j v^T x_{i,j}$$

Under this model, the (log) likelihood gradient is:

$$\frac{\partial \ell}{\partial w} = \frac{1}{2} \sum_i (y_i - \langle y_i \rangle_{\Theta; x}) x_i$$



# Taming the partition function

## Method II

Instead of approximating entire likelihood, try to estimate the problematic sum [Kumar et al., 2005]. Assume linear potentials ( $\Theta = (w, v)$ ):

$$\psi(y_i|x) = \log P(y_i|x) = y_i w^T x_i$$

$$\psi(y_i, y_j|x) = \log P(y_i, y_j|x) = y_i y_j v^T x_{i,j}$$

Under this model, the (log) likelihood gradient is:

$$\frac{\partial \ell}{\partial w} = \frac{1}{2} \sum_i (y_i - \langle y_i \rangle_{\Theta; x}) x_i$$

Need to estimate the expectation:

$$\langle y_i \rangle_{\Theta; x} = \sum_{y'} P(y'|x; \Theta) \cdot y'_i$$

# Taming the partition function

## Method II

- Marginal approximation:

# Taming the partition function

## Method II

- Marginal approximation:

Obtain estimated marginal probabilities per variable:

$$P_i(y_i|x; \Theta) = \sum_{y_1} \dots \sum_{y_{i-1}} \sum_{y_{i+1}} \dots \sum_{y_n} P(y_1, \dots, y_n|x; \Theta)$$

using loopy belief propagation.

# Taming the partition function

## Method II

- Marginal approximation:

Obtain estimated marginal probabilities per variable:

$$P_i(y_i|x; \Theta) = \sum_{y_1} \dots \sum_{y_{i-1}} \sum_{y_{i+1}} \dots \sum_{y_n} P(y_1, \dots, y_n|x; \Theta)$$

using loopy belief propagation. Then, plug marginals into expectation:

$$\langle y_i \rangle_{\Theta; x} = \sum_{y_i} y_i P_i(y_i|x; \Theta)$$

# Taming the partition function

## Method II

- Marginal approximation:

Obtain estimated marginal probabilities per variable:

$$P_i(y_i|x; \Theta) = \sum_{y_1} \dots \sum_{y_{i-1}} \sum_{y_{i+1}} \dots \sum_{y_n} P(y_1, \dots, y_n|x; \Theta)$$

using loopy belief propagation. Then, plug marginals into expectation:

$$\langle y_i \rangle_{\Theta; x} = \sum_{y_i} y_i P_i(y_i|x; \Theta)$$

- Saddle point approximation:

Let  $\hat{y} = \arg \max_y P(y|x; \Theta)$  be the most probable labeling.

# Taming the partition function

## Method II

- Marginal approximation:

Obtain estimated marginal probabilities per variable:

$$P_i(y_i|x; \Theta) = \sum_{y_1} \dots \sum_{y_{i-1}} \sum_{y_{i+1}} \dots \sum_{y_n} P(y_1, \dots, y_n|x; \Theta)$$

using loopy belief propagation. Then, plug marginals into expectation:

$$\langle y_i \rangle_{\Theta; x} = \sum_{y_i} y_i P_i(y_i|x; \Theta)$$

- Saddle point approximation:

Let  $\hat{y} = \arg \max_y P(y|x; \Theta)$  be the most probable labeling. Assume that the entire 'mass' of the partition function is concentrated at  $\hat{y}$ :

$$Z(\Theta; x) \approx \exp\left[\sum_i \psi(\hat{y}_i|x; \Theta) + \sum_i \sum_{j \in N(i)} \psi(\hat{y}_i, \hat{y}_j|x; \Theta)\right]$$

# Taming the partition function

## Method II

- Marginal approximation:

Obtain estimated marginal probabilities per variable:

$$P_i(y_i|x; \Theta) = \sum_{y_1} \dots \sum_{y_{i-1}} \sum_{y_{i+1}} \dots \sum_{y_n} P(y_1, \dots, y_n|x; \Theta)$$

using loopy belief propagation. Then, plug marginals into expectation:

$$\langle y_i \rangle_{\Theta; x} = \sum_{y_i} y_i P_i(y_i|x; \Theta)$$

- Saddle point approximation:

Let  $\hat{y} = \arg \max_y P(y|x; \Theta)$  be the most probable labeling. Assume that the entire 'mass' of the partition function is concentrated at  $\hat{y}$ :

$$Z(\Theta; x) \approx \exp\left[\sum_i \psi(\hat{y}_i|x; \Theta) + \sum_i \sum_{j \in N(i)} \psi(\hat{y}_i, \hat{y}_j|x; \Theta)\right]$$

Then,  $\langle y_i \rangle_{\Theta; x} \approx \hat{y}_i$ .

# Taming the partition function

## Learning/inference results

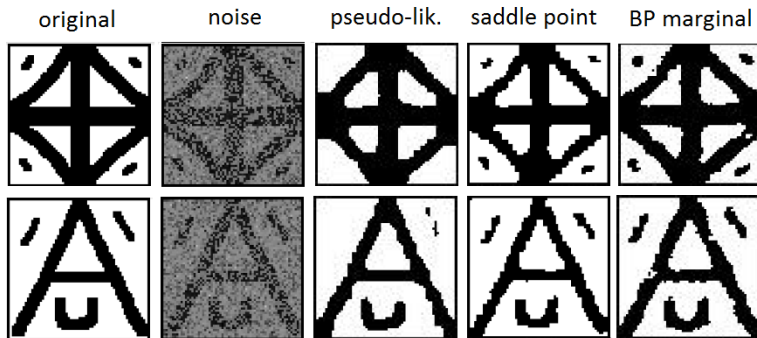


©Kumar et al., 2005



# Taming the partition function

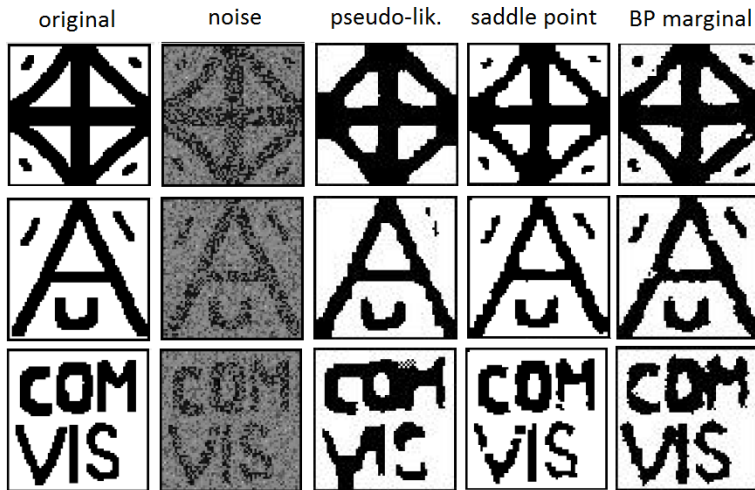
## Learning/inference results



©Kumar et al., 2005

# Taming the partition function

## Learning/inference results



©Kumar et al., 2005

# Max-margin learning

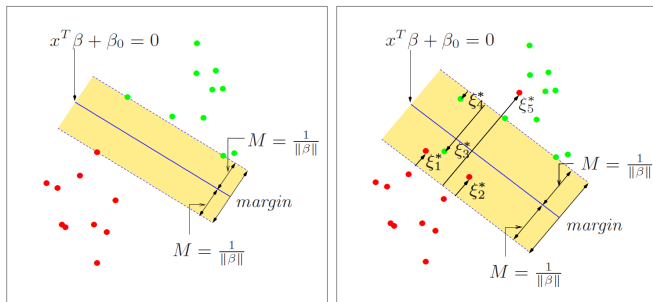
## Introduction

- Abandons probabilistic perspective
- Aims at maximizing energy *margin* between object labels
- Uses *linear* potentials, can be kernelized

# Max-margin learning

## Introduction

- Abandons probabilistic perspective
- Aims at maximizing energy *margin* between object labels
- Uses *linear* potentials, can be kernelized
- Generalizes ideas from support vector machine



©Hastie et al., 2008

# Max-margin learning

## Basic formulation

Let  $\psi(y_i) = w^T F(x_i, y_i)$ ,  $\psi(y_i, y_j) = v^T F(x_{i,j}, y_i, y_j)$  be the potentials.

# Max-margin learning

## Basic formulation

Let  $\psi(y_i) = w^T F(x_i, y_i)$ ,  $\psi(y_i, y_j) = v^T F(x_{i,j}, y_i, y_j)$  be the potentials.

- $\Theta = (w, v)$ : compound weight,  $F(x, y)$ : compound feature
- $\Delta(y, z) \geq 0$ : distance between two labelings  $y$  and  $z$
- Energy:  $E(y|x; \Theta) = \Theta^T F(x, y)$

# Max-margin learning

## Basic formulation

Let  $\psi(y_i) = w^T F(x_i, y_i)$ ,  $\psi(y_i, y_j) = v^T F(x_{i,j}, y_i, y_j)$  be the potentials.

- $\Theta = (w, v)$ : compound weight,  $F(x, y)$ : compound feature
- $\Delta(y, z) \geq 0$ : distance between two labelings  $y$  and  $z$
- Energy:  $E(y|x; \Theta) = \Theta^T F(x, y)$

Main idea: bias the weights  $\Theta$  so that the energy of training labels is lower than energy of any other configuration *by the maximum margin*

# Max-margin learning

## Basic formulation

Let  $\psi(y_i) = w^T F(x_i, y_i)$ ,  $\psi(y_i, y_j) = v^T F(x_{i,j}, y_i, y_j)$  be the potentials.

- $\Theta = (w, v)$ : compound weight,  $F(x, y)$ : compound feature
- $\Delta(y, z) \geq 0$ : distance between two labelings  $y$  and  $z$
- Energy:  $E(y|x; \Theta) = \Theta^T F(x, y)$

Main idea: bias the weights  $\Theta$  so that the energy of training labels is lower than energy of any other configuration *by the maximum margin*

Structured SVM [Taskar et al., 2003, Tsochantaridis et al., 2005]

maximize  $\gamma$

s.t.  $\|\Theta\| \leq 1$

$$\forall_{y' \neq y} \underbrace{\Theta^T F(x, y')}_{E(y')} - \underbrace{\Theta^T F(x, y)}_{E(y)} \geq \Delta(y, y') \cdot \gamma$$



# Max-margin learning

## Slack variable formulation

Training set may not be separable - no solution. Need to add slack variables:

# Max-margin learning

## Slack variable formulation

Training set may not be separable - no solution. Need to add slack variables:

$$\text{minimize } \frac{1}{2} \|\Theta\|_2^2 + \frac{C}{N} \sum_i \zeta_i$$

$$\text{s.t. } \zeta_i \geq 0$$

$$\forall_{x_i} \forall_{y'_i \neq y_i} \underbrace{\Theta^T F(x_i, y'_i)}_{E(y'_i)} - \underbrace{\Theta^T F(x_i, y_i)}_{E(y_i)} \geq \Delta(y_i, y'_i) - \zeta_i$$

# Max-margin learning

## Slack variable formulation

Training set may not be separable - no solution. Need to add slack variables:

$$\text{minimize } \frac{1}{2} \|\Theta\|_2^2 + \frac{C}{N} \sum_i \zeta_i$$


$$\text{s.t. } \zeta_i \geq 0$$

$$\forall_{x_i} \forall_{y'_i \neq y_i} \underbrace{\Theta^T F(x_i, y'_i)}_{E(y'_i)} - \underbrace{\Theta^T F(x_i, y_i)}_{E(y_i)} \geq \Delta(y_i, y'_i) - \zeta_i$$

Example: semantic segmentation of urban scenes [Volpi and Ferrari, 2015]

“In our setting, the potentials employed are linear with respect to the parameter vector.(...)To learn the CRF parameters **w** we adopt the margin rescaling variant of the SSVM(...)”

— Volpi and Ferrari



© Volpi and Ferrari

# Max-margin learning

## Main challenge

Constraints revisited:

$$\forall_{y' \neq y} \Theta^T F(x, y') - \Theta^T F(x, y) \geq \Delta(y, y') \cdot \gamma$$

# Max-margin learning

## Main challenge

Constraints revisited:

$$\forall_{y' \neq y} \Theta^T F(x, y') - \Theta^T F(x, y) \geq \Delta(y, y') \cdot \gamma$$

- One constraint for every possible label assignment
- Exponential in input data length ! Infeasible to add all constraints at once

# Max-margin learning

## Main challenge

Constraints revisited:

$$\forall_{y' \neq y} \Theta^T F(x, y') - \Theta^T F(x, y) \geq \Delta(y, y') \cdot \gamma$$

- One constraint for every possible label assignment
- Exponential in input data length ! Infeasible to add all constraints at once

Possible solution [Tsochantaridis et al., 2005]:

- Add constraints iteratively
- At every step, find most violated constraint
- Re-solve with augmented constraint set

# Max-margin learning

## Determining most violated constraint

Find the labeling minimizing the loss-augmented energy:

$$\hat{y} = \arg \max_{y'} E(y' | x, \underbrace{\Theta}_{\text{current est.}}) - \Delta(y', y)$$

# Max-margin learning

## Determining most violated constraint

Find the labeling minimizing the loss-augmented energy:

$$\hat{y} = \arg \max_{y'} E(y' | x, \underbrace{\Theta}_{\text{current est.}}) - \Delta(y', y)$$

Assume  $\Delta$  decomposes according to nodes:  $\Delta(y, y') = \sum_i \delta(y_i, y'_i)$ . Define augmented unary potentials:

$$\bar{\psi}(y'_i) = \psi(y'_i) - \delta(y_i, y'_i)$$



# Max-margin learning

## Determining most violated constraint

Find the labeling minimizing the loss-augmented energy:

$$\hat{y} = \arg \max_{y'} E(y' | x, \underbrace{\Theta}_{\text{current est.}}) - \Delta(y', y)$$

Assume  $\Delta$  decomposes according to nodes:  $\Delta(y, y') = \sum_i \delta(y_i, y'_i)$ . Define augmented unary potentials:

$$\bar{\psi}(y'_i) = \psi(y'_i) - \delta(y_i, y'_i)$$

Perform inference in new network with energy  $E(y'; \bar{\psi}(\cdot), \psi(\cdot, \cdot))$  using preferred method (QPBO, linear programming relaxation, simulated annealing, ...)



# Available software libraries

- UGM: Matlab code for undirected graphical models (<http://www.cs.ubc.ca/~schmidtm/Software/UGM.html>)
- SVM<sup>struct</sup>: Structured SVM API in Python, C++, Matlab ([http://www.cs.cornell.edu/people/tj/svm\\_light/svm\\_struct.html](http://www.cs.cornell.edu/people/tj/svm_light/svm_struct.html))
- PyStruct - Structured Learning in Python (<http://pystruct.github.io/>)
- OpenGM: a C++ template library for discrete factor graph models (<http://hciweb2.iwr.uni-heidelberg.de/opengm/>)

# Available software libraries

- UGM: Matlab code for undirected graphical models (<http://www.cs.ubc.ca/~schmidtm/Software/UGM.html>)
- SVM<sup>struct</sup>: Structured SVM API in Python, C++, Matlab ([http://www.cs.cornell.edu/people/tj/svm\\_light/svm\\_struct.html](http://www.cs.cornell.edu/people/tj/svm_light/svm_struct.html))
- PyStruct - Structured Learning in Python (<http://pystruct.github.io/>)
- OpenGM: a C++ template library for discrete factor graph models (<http://hciweb2.iwr.uni-heidelberg.de/opengm/>)



# Summary

- Numerous learning methods with varying degrees of complexity and capabilities

# Summary

- Numerous learning methods with varying degrees of complexity and capabilities
- Two main paradigms: probabilistic perspective, max-margin learning

# Summary

- Numerous learning methods with varying degrees of complexity and capabilities
- Two main paradigms: probabilistic perspective, max-margin learning
- Many specializations have been developed (chain graphs, trees, submodular potentials, . . .)

# Summary

- Numerous learning methods with varying degrees of complexity and capabilities
- Two main paradigms: probabilistic perspective, max-margin learning
- Many specializations have been developed (chain graphs, trees, submodular potentials, . . . )
- Direct methods may still be useful



Kumar, S., August, J., and Hebert, M. (2005).

*Exploiting Inference for Approximate Parameter Learning in Discriminative Fields: An Empirical Study*, pages 153–168.

Springer Berlin Heidelberg, Berlin, Heidelberg.



Niemeyer, J., Rottensteiner, F., and Soergel, U. (2014).

Contextual classification of lidar data and building object detection in urban areas.

*ISPRS Journal of Photogrammetry and Remote Sensing*, 87:152 – 165.



Schmidt, M., Murphy, K., Fung, G., and Rosales, R. (2008).

Structure learning in random fields for heart motion abnormality detection.

In *Computer Vision and Pattern Recognition, 2008. CVPR 2008. IEEE Conference on*, pages 1–8.



Taskar, B., Guestrin, C., and Koller, D. (2003).

Max-margin markov networks.

MIT Press.





Tsochantaridis, I., Joachims, T., Hofmann, T., and Altun, Y. (2005).  
Large margin methods for structured and interdependent output variables.  
*J. Mach. Learn. Res.*, 6:1453–1484.



Volpi, M. and Ferrari, V. (2015).  
Semantic segmentation of urban scenes by learning local class  
interactions.  
*In The IEEE Conference on Computer Vision and Pattern Recognition  
(CVPR) Workshops.*



Weinmann, M., Schmidt, A., Mallet, C., Hinz, S., Rottensteiner, F., and  
Jutzi, B. (2015).  
Contextual classification of point cloud data by exploiting individual 3d  
neighbourhoods.  
*ISPRS Annals of Photogrammetry, Remote Sensing and Spatial  
Information Sciences*, II-3/W4:271–278.

# The end

Thank you for your attention.